



Is Artificial Intelligence Really Undermining Democracy?

A Critical Appraisal of Mark Coeckelbergh's *Why AI Undermines Democracy* (2024)

Peter Kahl; independent researcher; first published 'free-range' 12 July 2025 on Substack

Abstract

Mark Coeckelbergh's recent book, *Why AI Undermines Democracy* (2024), argues compellingly that artificial intelligence (AI) intrinsically threatens democratic governance by facilitating epistemic manipulation, deepening knowledge asymmetries, and fostering political alienation. While acknowledging Coeckelbergh's valuable insights into democratic vulnerabilities exacerbated by digital technologies, this essay critiques his foundational assumption that AI inherently possesses political or epistemic agency. Drawing upon fiduciary epistemic theory—particularly elaborated in my works 'Epistemic Justice and Institutional Responsibility in Academia' (2015) and

“Directors’ Epistemic Duties and Fiduciary Openness” (2025)—I propose an alternative governance framework, termed ‘epistemocracy’, which explicitly locates epistemic and moral responsibility within human institutional actors rather than technological artefacts. Unlike Coeckelbergh’s ‘technodemocracy’, epistemocracy prioritises fiduciary transparency, institutional accountability, and epistemic pluralism as foundational to democratic governance of AI. Ultimately, the essay argues that attributing inherent political agency to AI obscures crucial institutional responsibilities and weakens democratic safeguards. Recognising AI as epistemically neutral and embedding fiduciary epistemic duties into governance structures provides clearer, more robust pathways for institutional reform, democratic accountability, and sustainable epistemic justice in AI governance.

Keywords

AI governance, artificial intelligence ethics, democracy, democratic accountability, democratic reform, digital humanism, epistemic agency, epistemic bubbles, epistemic justice, epistemic neutrality, epistemocracy, epistemic pluralism, fiduciary duties, fiduciary epistemic theory, governance frameworks, institutional accountability, institutional responsibility, institutional transparency, knowledge asymmetries, Mark Coeckelbergh, political epistemology, technological determinism, technocracy, technodemocracy

Table of Contents

Illustration Description and Analysis	4
1. Introduction	5
1.1 Context and Purpose.....	5
1.2 Summary of Coeckelbergh’s Central Arguments.....	5
1.3 Overview of Critical Perspective	5
1.4 Thesis Statement and Methodology	6
2. Coeckelbergh’s Key Contributions	6
2.1 AI, Technocracy, and Technodemocracy	6
2.2 Epistemic Bubbles, Knowledge Asymmetries, and Democracy	7
2.3 Digital Humanism and the Democratic Ideal.....	7
3. Critiquing the Premise: AI’s Political Agency	8
3.1 AI as a Politically Relational Technology?	8
3.2 The Problem of Moral and Epistemic Agency	8
3.3 Reasserting Epistemic Neutrality: AI as Instrument	9
4. Fiduciary Epistemic Framework: A Counterproposal.....	9
4.1 Institutional Epistemic Responsibility and Accountability.....	9
4.2 Fiduciary Duties in Epistemic Governance	10
4.3 Addressing Epistemic Injustice through Fiduciary Openness	11
5. Implications for Democratic Governance and AI Policy	11
5.1 Institutional Reforms over Technological Determinism	11
5.2 Democratic Accountability and AI Regulation.....	12
5.3 Enhancing Trust and Epistemic Pluralism through Institutional Transparency	12
6. Conclusion: Reframing the AI-Democracy Relationship	13
6.1 Summary of Arguments.....	13
6.2 Recommendations for Further Scholarship and Policy Development	14
Bibliography	14
Author Contact Information	15
Cite this work	15

Illustration Description and Analysis

The illustration accompanying this essay, generated using the DALL-E AI image-generation model in a style evocative of Roy Lichtenstein's iconic pop-art aesthetic, vividly encapsulates the central arguments presented within this critique. At the centre of the composition stands a stylised humanoid figure prominently labelled 'AI,' raising its right hand in a gesture symbolising neutrality and impartiality, evoking imagery typically associated with witnesses giving testimony. Its illuminated eyes suggest alertness, yet it remains expressionless, reinforcing the essay's assertion that AI possesses no inherent moral or political agency.

Flanking the AI figure are two barristers, each adorned in traditional British legal attire, including wigs and robes, representing competing philosophical perspectives in AI governance. On the left, one barrister firmly asserts, 'TOOLS REFLECT INTENT,' highlighting the essay's position that AI technologies merely amplify and reflect human purposes, values, and biases. Opposing him, the barrister on the right counters emphatically, 'GOVERNANCE SHAPES MEANING,' underscoring the essay's argument that the institutional governance framework fundamentally determines how AI impacts society.

The bold, dichromatic background—divided between vibrant crimson above and calm aquamarine below—symbolically captures the tension between alarmist narratives and reasoned governance approaches. This juxtaposition visually reinforces the essay's critical position against technological determinism, emphasising instead institutional accountability, fiduciary responsibility, and epistemic clarity. Collectively, the illustration powerfully underscores the central thesis: democratic governance of AI must be epistemically transparent and explicitly situated within human institutional responsibility, rather than attributing misleading political or epistemic agency to technology itself.

1. Introduction

1.1 Context and Purpose

In an era profoundly shaped by artificial intelligence (AI), societies find themselves at critical crossroads concerning democracy, governance, and epistemic justice. Technologies underpinning AI—such as machine learning algorithms, data analytics, and automated decision-making systems—have transformed how political institutions engage with citizens and manage public affairs. This transformation inevitably prompts questions about whether such technologies inherently bolster or threaten democratic values. Mark Coeckelbergh's recent book, *Why AI Undermines Democracy* (2024), contributes significantly to this discourse by contending that AI fundamentally erodes democratic principles.

This essay critically appraises Coeckelbergh's arguments, recognising his substantial contributions while robustly challenging his foundational assumptions. Its purpose is to clarify and refine our understanding of AI's role within democratic governance, with particular emphasis on epistemic accountability, fiduciary responsibility, and institutional agency.

Unless explicitly stated otherwise, all summaries and analyses of Coeckelbergh's arguments presented throughout this essay are based on his book *Why AI Undermines Democracy* (2024).

1.2 Summary of Coeckelbergh's Central Arguments

Coeckelbergh argues that AI poses intrinsic threats to democratic governance by enabling technocratic and authoritarian practices that exploit knowledge asymmetries and exacerbate epistemic divisions, such as echo chambers and epistemic bubbles. He asserts that AI is not merely a neutral instrument but a politically charged entity—'politically relational' technology—which influences and potentially manipulates human behaviour, thereby undermining the trust, inclusivity, and solidarity essential to democratic life. Coeckelbergh promotes 'technodemocracy' as an alternative to technocracy, urging a shift towards deliberative, participative, and republican conceptions of democracy. His vision involves embedding democratic ideals within AI technologies through regulatory measures and embracing a form of digital humanism informed by Renaissance and Enlightenment values, emphasising the common good, ethical responsibility, and active citizen participation.

1.3 Overview of Critical Perspective

While Coeckelbergh's critique insightfully highlights the challenges AI technologies present, it raises critical conceptual issues that warrant deeper scrutiny. Central to these issues is his claim that AI possesses inherent political agency or epistemic value. My critical perspective contends that this attribution risks obscuring the essential distinction between technological artefacts and human agency, potentially diluting institutional accountability and moral responsibility. I argue instead that AI is epistemically neutral and politically inert; harm or benefit arises entirely from the human agents, institutional frameworks, and epistemic governance practices surrounding its deployment. Building upon fiduciary epistemic theory elaborated in my works, notably 'Epistemic Justice and Institutional Responsibility in Academia' and "Directors' Epistemic Duties and Fiduciary Openness", I advocate for explicitly situating responsibility and accountability within institutions rather than within the technologies themselves. This epistemic-fiduciary approach seeks to clarify accountability, foster

transparency, and enhance democratic practices, thereby offering more practical avenues for mitigating epistemic injustice and democratic erosion.

1.4 Thesis Statement and Methodology

This essay argues that Coeckelbergh's characterisation of AI as inherently political or epistemically laden is conceptually and practically problematic, diverting focus from institutional actors and fiduciary responsibilities where true accountability must reside. My thesis maintains that AI, as a tool, possesses no inherent moral or political attributes; democratic erosion and epistemic injustice result exclusively from failures in institutional governance, fiduciary accountability, and epistemic transparency.

To substantiate this argument, the essay employs a critical analytical methodology. Firstly, it systematically reviews Coeckelbergh's core propositions, accurately summarising and contextualising them within the broader scholarship on AI and democracy. Secondly, through a detailed conceptual critique rooted in fiduciary epistemic theory, the essay deconstructs the attribution of agency to AI, highlighting conceptual gaps and practical limitations. Lastly, the essay outlines an alternative fiduciary epistemic framework, proposing institutional transparency and accountability as effective, actionable responses to democratic challenges posed by AI. Throughout, the methodology integrates analytical rigour with a nuanced approach to epistemic justice, maintaining clarity regarding theoretical claims and practical implications for democratic governance.

2. Coeckelbergh's Key Contributions

2.1 AI, Technocracy, and Technodemocracy

A central contribution of Mark Coeckelbergh's *Why AI Undermines Democracy* lies in his examination of technocracy and its complex relationship with artificial intelligence (AI). Coeckelbergh identifies technocracy as a governance structure predominantly led by technical experts, which increasingly relies on algorithmic decision-making, data-driven methodologies, and AI-driven processes. He argues convincingly that this form of governance tends to marginalise ordinary citizens by prioritising technical expertise and opaque algorithmic outcomes over inclusive democratic deliberation. As a result, technocracy reinforces power imbalances between technical elites—such as engineers, data scientists, and policy specialists—and the wider public, eroding core democratic values including citizen participation, transparent deliberation, and mutual trust.

To counteract this technocratic drift, Coeckelbergh introduces the concept of 'technodemocracy', which he describes as a governance model that proactively integrates republican and deliberative democratic principles directly into technological development and policy-making. He advocates a proactive stance, arguing that democratic values should not merely respond to technological innovations retrospectively but must be deliberately embedded within the initial conception, design, implementation, and governance frameworks of AI technologies. For Coeckelbergh, technodemocracy represents a fundamental shift from passive acceptance of technology's political implications towards active civic engagement, democratic oversight, and heightened institutional transparency.

2.2 Epistemic Bubbles, Knowledge Asymmetries, and Democracy

Another significant area addressed by Coeckelbergh is the epistemological impact of AI on democratic societies, particularly concerning the problems of epistemic bubbles, echo chambers, and knowledge asymmetries. Coeckelbergh argues that AI's reliance on algorithmically driven content delivery mechanisms intensifies epistemic fragmentation. By systematically presenting users with information tailored to reinforce existing biases or preferences, AI contributes to isolating individuals within ideological echo chambers, significantly weakening democratic discourse and civic cohesion. These epistemic bubbles insulate individuals and groups from exposure to diverse or opposing viewpoints, thereby restricting citizens' abilities to form well-informed, reflective, and autonomous political judgements.

Additionally, Coeckelbergh emphasises the critical challenge posed by knowledge asymmetry between AI experts and the general public. His analysis critiques the heavy dependence on specialised technical expertise characteristic of technocratic governance, highlighting the risks this presents to democratic accountability. Coeckelbergh warns that uncritical reliance on expert knowledge can lead societies towards epistemic authoritarianism, where interpretations by technical elites dominate public discourse, marginalising lay voices and undermining pluralistic democratic values. By questioning the conflation of expertise with absolute objectivity, Coeckelbergh advocates for greater epistemic inclusivity, asserting that the democratisation of knowledge production and dissemination is essential for fostering genuine democratic governance around AI.

2.3 Digital Humanism and the Democratic Ideal

Coeckelbergh's third notable contribution is his advocacy of 'digital humanism', a philosophical approach aimed explicitly at aligning technological development with democratic and humanist principles. Drawing intellectual inspiration from Renaissance humanism and Enlightenment values, digital humanism proposes that AI and related technologies should primarily serve ethical, democratic, and socially beneficial ends, as opposed to narrow commercial interests or instrumental objectives alone. Coeckelbergh suggests that this shift in orientation could substantially mitigate the risk of AI being exploited in ways that harm democratic values, social justice, and community solidarity.

Central to digital humanism is the explicit integration of fundamental democratic and ethical values—such as inclusivity, fairness, solidarity, mutual respect, and trust—directly within technological infrastructures and AI design practices. Coeckelbergh argues persuasively that failing to adopt such a humanist stance risks deepening socioeconomic inequalities, fostering isolation and alienation, and potentially triggering severe democratic erosion or collapse. He insists that maintaining and enhancing democratic governance necessitates that technological development explicitly addresses deeper moral considerations about human flourishing, collective well-being, and the common good. This represents an essential ethical dimension in contemporary debates around technology governance and democratic sustainability.

Summary

In sum, Coeckelbergh's work contributes substantially to contemporary discourses on AI and democracy by explicitly addressing the risks posed by technocratic governance, epistemic fragmentation, and advocating for morally reflective and democratically accountable technological frameworks. While his critiques and proposals significantly enrich scholarly and policy discussions, subsequent chapters of this essay will critically evaluate and challenge his underlying assumptions regarding the inherent political and epistemic attributes of AI. Instead, the essay argues, the primary focus should be placed upon clarifying institutional accountability,

fiduciary duties, and the epistemic responsibilities of human actors who govern AI systems and technological infrastructures.

3. Critiquing the Premise: AI's Political Agency

3.1 AI as a Politically Relational Technology?

Coeckelbergh's foundational claim—that artificial intelligence (AI) is inherently political or 'politically relational'—demands careful scrutiny. At the heart of this claim lies the suggestion that AI technologies embody specific political and epistemic attributes, actively influencing democratic processes and shaping human interactions within society. Coeckelbergh asserts that AI is not simply an instrument of governance but inherently shapes political relations, thereby attributing to it qualities of intentional agency typically reserved for human actors. While this view insightfully foregrounds the profound implications AI can have within political spheres, it conflates the technological artefact with the socio-political environment in which the artefact is deployed.

Coeckelbergh's framing potentially obscures critical distinctions between the technological systems themselves and the contexts of their implementation. To suggest AI is inherently 'politically relational' risks inadvertently shifting accountability away from clearly identifiable human decision-makers, institutional governance frameworks, and fiduciary actors. Furthermore, such characterisation arguably reflects a conceptual misunderstanding of AI's technical nature. Although AI systems undeniably affect democratic deliberation through algorithmically mediated content, these effects are contingent on human-designed parameters, training datasets, and deliberate institutional choices, rather than intrinsic political qualities embedded in the technology itself. Recognising this distinction is crucial for accurate analyses of democratic risks and for identifying appropriate policy responses.

3.2 The Problem of Moral and Epistemic Agency

Assigning inherent political agency to AI technologies raises significant philosophical concerns, particularly regarding moral and epistemic accountability. Coeckelbergh's view implies that AI, by virtue of algorithmic mediation and automated decision-making, may share moral responsibility or epistemic agency previously reserved exclusively for human actors. However, such attribution is deeply problematic. Moral agency implies conscious intention, reflective deliberation, and accountability—characteristics fundamentally absent from algorithmic systems, irrespective of their complexity. Similarly, epistemic agency involves intentional cognitive engagement and judgement, qualities AI systems categorically lack.

Ascribing moral and epistemic agency to AI inadvertently dilutes accountability, creating conceptual confusion regarding where moral responsibility resides. If political harm or epistemic injustice arises, it must be clearly attributable to identifiable agents capable of reflective intention and ethical judgement. By implicitly attributing agency to AI, Coeckelbergh risks obscuring the responsibilities of institutional actors, human designers, and governance structures that control AI's parameters and deployment. Crucially, moral and epistemic accountability must remain situated within these human contexts to ensure robust democratic oversight, transparency, and fiduciary responsibility.

3.3 Reasserting Epistemic Neutrality: AI as Instrument

In contrast to Coeckelbergh’s framing, I propose that AI should be understood as epistemically neutral—an instrumental tool whose political and epistemic impacts derive entirely from human design, intent, and institutional contexts. Drawing upon fiduciary epistemic theory articulated in my prior works, notably ‘Epistemic Justice and Institutional Responsibility in Academia’ (2025), “Directors’ Epistemic Duties and Fiduciary Openness” (2025), and ‘Epistemocracy in Higher Education’ (2025), I argue explicitly for a shift towards epistemocracy in AI governance. Unlike Coeckelbergh’s concept of technodemocracy, epistemocracy recognises that the governance of AI must fundamentally be guided by robust epistemic principles and fiduciary duties, rather than superficially embedding democratic ideals into technological structures without sufficient consideration of institutional epistemic accountability.

As a deep-technology expert, I maintain that AI systems fundamentally operate through mathematical algorithms and computational processes devoid of intrinsic moral or political content. Algorithms perform pattern recognition, prediction, and decision-making based strictly on data inputs and design parameters set by human actors. Hence, biases, political preferences, or epistemic injustices reflected in AI outputs are direct manifestations of decisions made by human programmers, designers, policymakers, or institutional actors.

An epistemic framework explicitly prioritises fiduciary epistemic governance, placing accountability squarely within institutional structures and fiduciary actors—individuals or entities entrusted with ethical oversight and epistemic responsibility. Fiduciary duties under epistemocracy explicitly demand rigorous transparency, inclusivity, and fairness in AI design and implementation. Institutional actors are thereby held directly accountable for epistemic injustices or democratic erosion associated with AI systems, ensuring clear lines of moral and epistemic responsibility.

Unlike technodemocracy, which risks obscuring the epistemic foundations of AI governance, epistemocracy explicitly foregrounds epistemic openness and institutional transparency. By clearly defining fiduciary epistemic duties, epistemocracy ensures actionable accountability mechanisms, promotes epistemic pluralism, and addresses democratic vulnerabilities proactively. Epistemocracy thus provides a more rigorous, ethically grounded alternative for safeguarding democratic values and epistemic justice, highlighting transparent human accountability rather than ambiguous technological agency.

In conclusion, while Coeckelbergh’s insights regarding AI’s democratic implications are undeniably valuable, attributing inherent political or epistemic agency to AI itself constitutes a conceptual and practical misstep. Clarifying the epistemic neutrality of AI through the explicit adoption of epistemocracy realigns democratic accountability with human institutional actors, providing clearer, more effective paths towards epistemically just and democratically robust governance of AI technologies.

4. Fiduciary Epistemic Framework: A Counterproposal

4.1 Institutional Epistemic Responsibility and Accountability

In addressing the epistemic and democratic challenges posed by artificial intelligence (AI), it is essential to articulate a clear institutional framework centred explicitly on epistemic responsibility and accountability. The preceding critique demonstrates that assigning inherent political or epistemic agency to AI itself inadvertently obscures the pivotal roles and responsibilities of human actors and institutional decision-makers. To redress this,

I propose adopting an explicit fiduciary epistemic framework, grounded firmly in institutional accountability and clearly delineated epistemic responsibilities.

Central to this approach is the concept of epistemic responsibility, defined as institutions' obligation to ensure the integrity, openness, and plurality of knowledge produced and disseminated through AI-driven platforms. Institutional epistemic responsibility includes actively addressing biases, promoting transparency, and ensuring fair representation of diverse epistemic communities. Accountability under this framework remains unambiguously located within human institutions and fiduciaries—such as governmental bodies, regulatory authorities, corporate directors, and public sector leaders—who govern AI's design, implementation, and oversight.

Importantly, institutions must ensure rigorous transparency in algorithmic decision-making processes, openly communicate their epistemic standards, and be held accountable for failures leading to epistemic harms or injustices. Institutions must proactively disclose how AI is designed, trained, and deployed, facilitating public understanding and informed democratic engagement. Only through clearly articulated institutional accountability can democratic values such as fairness, transparency, and inclusivity be effectively realised.

4.2 Fiduciary Duties in Epistemic Governance

A central pillar of the fiduciary epistemic framework involves applying fiduciary principles explicitly to the epistemic governance of AI. Drawing insights from my prior research in "Directors' Epistemic Duties and Fiduciary Openness" (2025), fiduciary duties—traditionally comprising duties of care, loyalty, and good faith—are here expanded to encompass explicit epistemic obligations. Within democratic governance contexts, fiduciaries must uphold duties to maintain epistemic fairness, openness, and pluralism in their stewardship of AI technologies and knowledge resources.

Epistemic fiduciaries have specific obligations, such as the duty of epistemic openness, which demands proactive transparency in AI's epistemic processes and outcomes. This duty includes transparently documenting algorithmic criteria, decision-making rationales, and data sourcing practices. Fiduciaries are also obliged to uphold duties of epistemic fairness and inclusivity, which require them to actively seek, incorporate, and represent diverse epistemic perspectives, particularly those historically marginalised or silenced.

This fiduciary epistemic approach to AI governance aligns closely with my earlier work, 'Epistemocracy in Higher Education' (2025). In that work, I introduced epistemocracy as a governance framework specifically for universities, explicitly foregrounding fiduciary accountability, epistemic openness, and transparent institutional responsibility to safeguard epistemic justice. I argued there that institutional fiduciaries—such as university administrators and academic governance bodies—bear explicit duties of epistemic transparency and inclusive knowledge governance.

By extension, applying epistemocracy to AI governance similarly mandates that institutions responsible for deploying AI must actively fulfil fiduciary epistemic obligations. These include clearly articulated duties of epistemic fairness, openness, transparency, and inclusivity in the development, deployment, and regulation of AI technologies. Contrasting epistemocracy explicitly with Coeckelbergh's technodemocracy further underscores epistemocracy's rigorous focus on fiduciary institutional accountability, rather than superficially embedding democratic ideals into technological artefacts themselves. Thus, epistemocracy ensures robust accountability structures clearly situated within human institutional actors rather than ambiguous technological processes.

In this fiduciary context, institutional leaders and governance bodies become explicitly accountable for epistemic injustices stemming from negligent or biased AI deployment. Rather than attributing moral or epistemic agency to AI, fiduciary duties unambiguously direct institutional attention towards human accountability, encouraging proactive oversight and ethical vigilance. By defining and codifying fiduciary epistemic duties, this framework ensures robust institutional responsibility, clear accountability pathways, and democratic legitimacy.

4.3 Addressing Epistemic Injustice through Fiduciary Openness

A critical benefit of adopting the fiduciary epistemic framework lies in its potential to effectively address epistemic injustices arising from AI governance practices. Epistemic injustice, broadly defined as unfair exclusion, misrepresentation, or marginalisation of certain epistemic groups or perspectives, remains a significant challenge exacerbated by algorithmically mediated knowledge dissemination. Fiduciary openness offers a practical, ethically robust mechanism for redressing these injustices by mandating institutional transparency, fairness, and active epistemic inclusivity.

Fiduciary openness obliges institutions to disclose openly how AI algorithms determine what knowledge is promoted, silenced, or marginalised. This openness extends beyond mere transparency, requiring institutional fiduciaries to demonstrate active responsiveness to diverse epistemic inputs and concerns. Fiduciaries must establish clear, accountable mechanisms for epistemic complaints or grievances, actively correcting algorithmic biases and fostering greater epistemic justice through inclusive decision-making practices.

Moreover, fiduciary openness empowers marginalised communities by giving them explicit institutional channels to contest epistemic exclusion, bias, or harm. Fiduciaries bear explicit responsibility for ensuring that epistemic injustices are actively identified, acknowledged, and rectified, rather than hidden within opaque algorithmic processes. By embedding fiduciary openness within democratic AI governance frameworks, institutions actively cultivate epistemic pluralism, trust, and democratic legitimacy, strengthening civic solidarity and participation.

In conclusion, the fiduciary epistemic framework presents a viable, ethically grounded alternative to attributing inherent political or epistemic agency to AI itself. Clearly defining institutional epistemic responsibilities, fiduciary duties, and accountability structures enables targeted, actionable responses to epistemic injustices and democratic erosion. By situating moral and epistemic agency explicitly within human institutional actors, fiduciary epistemic governance provides a robust, practical foundation for democratic resilience in the age of AI.

5. Implications for Democratic Governance and AI Policy

5.1 Institutional Reforms over Technological Determinism

Building upon the fiduciary epistemic framework developed in the preceding chapters, it becomes evident that safeguarding democracy in the age of artificial intelligence (AI) demands explicit institutional reforms rather than relying on narratives of technological determinism. Technological determinism, implicitly evident in Coeckelbergh's approach, suggests an inevitable erosion of democratic values due to intrinsic properties of AI technologies. However, as argued throughout this essay, this approach mistakenly attributes political and epistemic agency to technological artefacts, inadvertently shifting attention away from the urgent and actionable reforms necessary within human institutional governance structures.

Institutional reforms must therefore explicitly embed fiduciary epistemic duties within regulatory authorities, governmental agencies, educational institutions, and private entities responsible for overseeing and implementing AI. Institutions should adopt governance structures explicitly oriented toward epistemic transparency, comprehensive algorithmic accountability, and clear democratic oversight. Practical reforms might include establishing independent oversight committees for AI ethics and accountability, creating statutory obligations mandating algorithmic transparency, and implementing formal institutional mechanisms to address epistemic grievances and injustices. Crucially, these reforms explicitly restore responsibility to human fiduciaries and institutional decision-makers, empowering robust democratic engagement rather than passively accepting the impacts of AI as technologically predetermined.

5.2 Democratic Accountability and AI Regulation

Central to meaningful institutional reform is the enhancement of democratic accountability within AI regulation frameworks. AI technologies, characterised by algorithmic complexity and inherent opacity, present distinct challenges to traditional mechanisms of democratic accountability. Overcoming these challenges requires regulatory frameworks that explicitly embed fiduciary epistemic duties, ensuring governance structures remain transparently accountable and inclusive. Regulatory authorities should therefore be legally obligated to disclose the epistemic criteria underpinning their decisions, proactively document algorithmic methodologies, and maintain public accountability for all AI-related policy decisions.

Concrete regulatory measures could include statutory obligations for regular public reporting on AI system impacts, mandatory disclosures concerning algorithmic training data and decision-making criteria, and legislative frameworks clearly establishing fiduciary epistemic duties for entities designing, deploying, or managing AI systems in public or private governance contexts. Additionally, regulatory frameworks should provide robust oversight procedures—such as citizen juries, independent regulatory audits, and publicly accessible grievance mechanisms—to structurally embed and practically enforce democratic accountability.

Implementing these enhanced accountability mechanisms ensures AI governance explicitly promotes transparency, trustworthiness, and epistemic justice, significantly strengthening democratic legitimacy. Crucially, regulatory practices grounded in fiduciary epistemic duties restore human institutional accountability, clearly delineating democratic oversight responsibilities rather than deferring them to ambiguous technological agency.

5.3 Enhancing Trust and Epistemic Pluralism through Institutional Transparency

Finally, adopting the fiduciary epistemic framework demands institutional transparency as a central mechanism for enhancing public trust and epistemic pluralism. Transparency in AI governance involves not merely openness in decision-making and technical processes but the active inclusion and representation of diverse epistemic communities within institutional deliberations and regulatory practices. Transparent disclosure practices ensure citizens clearly understand AI's precise impacts on democratic deliberation, public policy decisions, and epistemic representation, thus reinforcing public trust, civic solidarity, and institutional legitimacy.

Operationalising institutional transparency involves mandating open disclosures of algorithmic logic, data sources, decision-making criteria, and epistemic assumptions underlying AI systems. These transparency initiatives ensure AI technologies do not become perceived as opaque or immune to democratic scrutiny. Institutional transparency further fosters epistemic pluralism by explicitly requiring the representation of

diverse stakeholder perspectives and epistemic communities in AI governance deliberations, proactively addressing epistemic exclusion or marginalisation. Inclusive deliberative forums, public consultations, and accessible oversight channels enable diverse epistemic voices to participate actively and meaningfully in shaping AI-related policy decisions.

Transparency combined with epistemic inclusivity enables institutions to actively counteract epistemic fragmentation, reducing epistemic bubbles and echo chambers reinforced by algorithmically mediated information environments. Through robust, transparent governance mechanisms, citizens gain confidence in the fairness, integrity, and legitimacy of AI-informed democratic processes, thereby cultivating epistemic pluralism, enhancing institutional trust, and strengthening democratic resilience.

Conclusion

In sum, adopting the fiduciary epistemic framework outlined in this essay necessitates substantial institutional reform, strengthened democratic accountability, and explicit transparency mechanisms structured around fiduciary epistemic duties. Rejecting technological determinism, epistemocracy emphasises explicitly placing human accountability and fiduciary responsibility at the forefront of democratic governance. By clearly situating epistemic responsibilities within human institutional frameworks rather than technological artefacts, democratic societies can effectively mitigate AI's risks and actively harness its potential for epistemic justice, robust trust, and inclusive democratic engagement.

6. Conclusion: Reframing the AI-Democracy Relationship

6.1 Summary of Arguments

This essay has critically appraised Mark Coeckelbergh's *Why AI Undermines Democracy*, acknowledging its substantial contributions while robustly challenging its fundamental premise—the attribution of inherent political and epistemic agency to artificial intelligence (AI). Although Coeckelbergh provides valuable insights into democratic vulnerabilities exacerbated by technocratic governance, epistemic fragmentation, and knowledge asymmetries, his characterisation of AI as inherently 'politically relational' inadvertently diverts crucial attention away from human accountability, institutional responsibilities, and fiduciary obligations.

In contrast, I have argued explicitly for an epistemocratic perspective in AI governance, proposing that AI should be understood as epistemically neutral and politically inert. Under this epistemocratic framework—developed from fiduciary epistemic theory articulated in my earlier works, notably 'Epistemic Justice and Institutional Responsibility in Academia' (2025), "Directors' Epistemic Duties and Fiduciary Openness" (2025), and 'Epistemocracy in Higher Education' (2025)—AI's democratic impacts are understood as wholly contingent upon clearly delineated human governance, rigorous institutional oversight, and explicit fiduciary epistemic responsibilities. Central to epistemocracy is the principle that democratic governance structures must explicitly prioritise epistemic transparency, institutional accountability, and fiduciary responsibility, rather than relying on superficial embedding of democratic ideals within technological artefacts.

The essay further demonstrates that reframing AI's democratic role from technological determinism towards explicit human and institutional accountability significantly enhances democratic legitimacy, public trust, epistemic pluralism, and resilience. Epistemocracy thus offers not merely theoretical clarity, but also practical

governance pathways for mitigating AI's risks, safeguarding epistemic justice, and ensuring robust democratic practices in an increasingly complex technological landscape.

6.2 Recommendations for Further Scholarship and Policy Development

The epistemic approach articulated in this essay clearly indicates several important avenues for further scholarly investigation and practical policy-making. Future research should continue refining fiduciary epistemic governance theory, empirically exploring concrete policy applications, and undertaking comparative institutional analyses. Further scholarship must actively integrate multidisciplinary expertise, including philosophical, ethical, legal, technological, and sociopolitical dimensions, to fully operationalise epistemic governance in diverse governance contexts.

Additionally, epistemic governance highlights an important shift in contemporary narratives surrounding AI's societal impacts. While public discourse frequently emphasises fears about AI displacing human labour, the analysis presented here suggests precisely the opposite scenario. Effective governance of AI technologies will inevitably necessitate increased demand for specialists—philosophers, multidisciplinary researchers, ethicists, technologists, and legal experts—capable of rigorous epistemic oversight and responsible governance. Rather than displacing opportunities, epistemic governance of AI will generate significant new demands for intellectual, practical, and ethical expertise across numerous epistemic communities.

Recognising this shift further underscores our collective epistemic duty, not merely to democratic governance but also explicitly to knowledge itself, as articulated in my ongoing research including my paper on epistemic transposition. Institutions, scholars, and policymakers must actively embrace epistemic responsibilities, explicitly safeguarding democratic values alongside the integrity, openness, and inclusivity of our collective knowledge ecosystems.

In conclusion, reframing the AI-democracy relationship through epistemic governance provides a superior philosophical and practical foundation compared to Coeckelbergh's technodemocracy. Epistemic governance explicitly prioritises human epistemic accountability, fiduciary transparency, and clearly articulated institutional responsibilities in AI governance. By proactively embedding fiduciary epistemic principles within democratic institutional frameworks, societies can effectively navigate AI's profound democratic implications, ensuring robust accountability, transparent epistemic governance, and sustainable democratic resilience.



Bibliography

Coeckelbergh M, *Why AI Undermines Democracy* (Polity Press 2024)

Kahl P, 'Directors' Epistemic Duties and Fiduciary Openness' (2025) available at <<https://pkahl.substack.com/p/directors-epistemic-duties-and-fiduciary>> accessed 8 July 2025

—— 'Epistemic Justice and Institutional Responsibility in Academia' (2025) available at <<https://pkahl.substack.com/p/epistemic-justice-and-institutional>> accessed 8 July 2025

—— 'Epistemic governance in Higher Education: A Proposal for Fiduciary and Epistemic Accountability in the University' (2025) available at <<https://pkahl.substack.com/p/epistemic-governance-higher-education-fiduciary-epistemic-accountability>> accessed 8 July 2025

—— ‘Epistemic Humility and the Transposition of Ethical Duties into Epistemic Duties’ (Github, 2025) available at <<https://github.com/Peter-Kahl/Epistemic-Humility-and-the-Transposition-of-Ethical-Duties-into-Epistemic-Duties>> accessed 8 July 2025



Author Contact Information

Author: Peter Kahl, independent researcher

ORCID: <<https://orcid.org/0009-0003-1616-4843>>

Email: <peter.kahl@juris.vc>

LinkedIn: <<https://www.linkedin.com/in/peter-kahl-law/>>

GitHub: <<https://github.com/Peter-Kahl>>

PhilPapers: <<https://philpeople.org/profiles/peter-kahl>>

Google Scholar: <<https://scholar.google.com/citations?hl=en&user=z-yfRRYAAAAJ>>

Blog: <<https://pkahl.substack.com/>>

Correspondence regarding this paper is welcome.



Cite this work

Peter Kahl, ‘Is Artificial Intelligence Really Undermining Democracy? A Critical Appraisal of Mark Coeckelbergh’s Why AI Undermines Democracy (2024)’ (2025) available at <<https://github.com/Peter-Kahl/Is-Artificial-Intelligence-Really-Undermining-Democracy>>